

Exercises – Deep Learning

For this set of exercises we will be using Keras [<http://keras.io/>], a high-level neural networks library, written in Python and capable of running on top of either TensorFlow or Theano. To install Keras, use PyPI with: `sudo pip install keras`

The default backend library in Keras is TensorFlow but it requires an additional installation. For these exercises, change the backend to Theano [<http://keras.io/backend/>].

1. Neural Network

Open the `nn.py` and edit only the `main()` function.

- a) Load the MNIST dataset that is included in the Keras library using the `loadMNISTdataset()` function

```
x_train, y_train, x_test, y_test = loadMNISTdataset()
```

- b) Create a multilayer perceptron model and interpret its architecture using the Keras documentation [<http://keras.io/layers/core/>]:

```
model = Sequential()
model.add(Dense(64, input_dim=28*28, init='uniform', activation='sigmoid'))
model.add(Dense(10, init='uniform', activation='sigmoid'))
model.compile(loss='mse', optimizer='adam', metrics=['accuracy'])
```

- c) Train the model

```
history = model.fit(x_train, y_train, batch_size=64, nb_epoch=3, verbose=1,
                    validation_split=0.1)
```

Plot the accuracy of the training process using the function `plotTrainingHistory()`. Try different number of epochs (3, 10, 20) and interpret the results.

- d) Test the model

```
score = model.evaluate(x_test, y_test)
print("\nTest accuracy: %0.05f" % score[1])
```

Show the images that are erroneously classified using the function `showErrors()`

2. Convolutional Neural Network

Open the `cnn.py` and edit only the `main()` function.

- a) Load the MNIST dataset that is included in the Keras library using the `loadMNISTdataset()` function

```
x_train, y_train, x_test, y_test, input_shape = loadMNISTdataset()
```

- b) Create a multilayer perceptron model and interpret its architecture using the Keras documentation [<http://keras.io/layers/core/>]:

```
nb_filters = 32
pool_size = (2, 2)
kernel_size = (3, 3)
model = Sequential()
model.add(Convolution2D(nb_filters, kernel_size[0], kernel_size[1],
                        border_mode='valid',
                        input_shape=input_shape))
```

```

model.add(Activation('relu'))
model.add(Convolution2D(nb_filters, kernel_size[0], kernel_size[1]))
model.add(Activation('relu'))
model.add(MaxPooling2D(pool_size=pool_size))
model.add(Dropout(0.25))
model.add(Flatten())
model.add(Dense(128))
model.add(Activation('relu'))
model.add(Dropout(0.5))
model.add(Dense(10))
model.add(Activation('softmax'))
model.compile(loss='categorical_crossentropy',
              optimizer='adadelta',
              metrics=['accuracy'])

```

c) Train the model (it may take a while...)

```

history = model.fit(x_train, y_train, batch_size=64, nb_epoch=3, verbose=1,
                  validation_split=0.1)

```

d) Test the model

```

score = model.evaluate(x_test, y_test)
print("\nTest accuracy: %0.05f" % score[1])

```

- e) Since training a complex model can take a very long time to train, they can be saved and loaded later. Explore how this can be done [<https://keras.io/models/about-keras-models/>]
- f) It is also possible to use very complex models with pre-trained weights, such as VGG and ResNet. Explore how this can be done [<http://keras.io/applications/>]